

# Distributed System Services

Mirek Riedewald



This work is licensed under the Creative Commons Attribution 4.0 International License. To view a copy of this license, visit <http://creativecommons.org/licenses/by/4.0/>.

# Key Learning Goals

- How do all the services discussed in this module achieve consensus in distributed decision-making?

# Overview

- Scalable distributed data processing relies on sophisticated services and infrastructure in the distributed system. Here we focus on those most essential for understanding execution and performance of a MapReduce or Spark job:
  - File system
  - Resource management
  - Application management
- Other services, including for consensus, database, and key-value store are either not covered in this course, or are covered in a later module.